

Book of Magne

INF222 Crashcourse 2023

Sander Wiig

INF222 Crashcourse for v2023.



UNIVERSITY OF BERGEN
Faculty of Mathematics and Natural Sciences

Institute for Informatics
Bergen Language Design Laboratory
The University of Bergen
Bergen, Norway
May 14, 2023

Contents

1	What is a language	6
1.1	What is a programming language	6
1.2	Meta Programming	7
1.3	Sum of products	8
2	Anatomy of an Interpreter	10
2.1	Phases of an interpreter	10
2.2	ASTs	12
2.3	Static Analysis	13
2.3.1	Type Checking	13
2.3.2	Wellformedness	16
3	Store, State, and Storage	17
3.1	State	17
3.1.1	Store	17
3.1.2	Variables and Enviroment	17
3.2	Scoping	17
3.3	Data Structures in Memory	19
3.3.1	Arrays	19
3.3.2	Matrixes and Multidimensional Arrays	20
3.3.3	Records, Structs	22
4	Procedures	24
4.1	Procedure Declarations	24
4.1.1	Parameters & Local Variables	25
4.1.2	Performing a Procedure	25
4.2	Executing a Procedure Call	25
4.3	Arg Passing	26
4.3.1	Copy Semantics	26
4.3.2	Reference Semantics	28
5	Signatures	30
5.1	A little bit of Theory	30
5.1.1	Signatures	30
5.1.2	Algebras	30
5.2	Implementing Algebraic Specifications	31
5.2.1	ADTs	33
5.3	Generic Programming	35

5.3.1 Concepts	35
A Assertions	36
A.1 Types of Assertions	36
B Language Standards	37
B.1 Software Engineering Implications of Languages	37
B.2 Reading specifications	37
B.2.1 Backus-Naur Form	37
Glossary	39
Bibliography	41

Preface

The goal of this script is to provide an overview of the course INF222 at the University of Bergen. It is by no means a comprehensive guide to the course, but rather a supplement to the lectures and exercises.

This was all written throughout a weekend and there is probably a whole lot wrong with the text and the code, so if you find any errors, please let me know by submitting a ticket or pull request on GitHub.

The Glossary is especially lacking. Due to the time constraints, I was not able to update all terms in the glossary, so if you find a term that is not defined, or seems wrong, please let me know.

Contributors

The following people have contributed in some way to this book:

- **Magne Haveraaen** for going over the course outline, and for teaching me what I know.
- **Anya Bagge** for her excellent INF225 notes, and for being a great teacher[2].
- **Jørn Lode** for his amazing figures that illustrate how states work.
- **Ralf Lämmel** for his book on Software Language Engineering, which much of this book is based on[1].



Figure 1: <https://github.com/Swi005/Book-of-Magne/tree/v2023>

Chapter 1

What is a language

1.1 What is a programming language

A programming language

- is an artificial language(i.e made by us humans on purpose)
- used to tell machines what to do

More formally a programming language is a set of rules that converts some input, like strings, into instructions that the computer can follow. This is of course a very general description and it, therefore, follows that there are many different types of programming languages. IT therefore should come as little surprise that we group languages by features and properties.

Types of languages

There are many ways of grouping languages. They can be grouped by Purpose, typing, paradigm, Generality vs. Specificity, and many more. For now, we're going to group them by paradigm, and Generality vs. Specificity.

Generality vs. Specificity

Languages are usually grouped into two categories when based on their specificity.

- DSL
- GPL

Domain Specific Languages are as the name suggests languages with a "specific" domain. DSLs usually have limited scope and use. Examples are JSON and SQL. A Domain-Specific Language is a programming language with a higher level of abstraction optimized for a specific class of problems. Optimized for a certain problem/domain. DSLs can be further subdivided into external DSL(separate programming languages), and internal DLS(language-like interface as a library.)

The languages we have implemented can also be categorized as either a DSL or a GPL.

Basic Typed Language and **Basic Imperative Programing Language** are DSLs because they have small ASTs, specific domain(teaching languages), and are not user extendable

Procedural Imperative Programming Language however can be categorized as a GPL

because, while it shares the domain and purpose of BIPL and BTL it is user extendable, and this is enough to push it into the category of GPLs.

General Purpose Languages however are more general and can be used to solve many different problems in many different situations. These languages have a wide array of uses and are usually what we think of when we hear the words programming language. Examples of GPLs are Java and Haskell.

Characteristic	DSL	GPL
Domain	Small and well-defined domain	Generality, many use cases
Size	Small ASTs	Large ASTs, often user extensible
Lifespan	As long as their domain	years to decades
Extensibility	Usually not extensible by users	Provides mechanisms for extensibility

Figure 1.1: Comparison between GPLs and DSLs

Syntax and Semantics

All programming languages have two parts; the **Syntax**, and the **Semantics**.

Syntax is the study of *structure*, just as semantics is the study of *meaning*. Or in other words, the syntax tells us *how* to write legal programs, and the semantics tells us *what* those programs do.

1.2 Meta Programming

One of the harder things in the course is **Meta-Programming**. INF222 is usually the first time you've encountered meta-programming and it can be hard a hard concept to grasp. Meta-programming is programming *about* programming. More properly meta-programs treat other programs as data. Interpreters and compilers are one example of meta-programs.

When talking about a language in the context of a compiler it is often difficult to distinguish between the language itself and the language that the compile is written in. We call the language that the compiler is written in the **Meta-Language**, and the language that the compiler is compiling is the **Object Language**.

For example in the BTL interpreter, the object language is BTL, and the meta-language is Haskell.

When you see a data structure like **Basic Typed Language** or **Basic Imperative Programming Language** in Haskell it represents a program. *TODO: Talk about meta languages, Object languages, and classify BTL, BIPL, and PIPL as DSL or GPL*

1.3 Sum of products

You may have encountered the term **Sum of Products**, let's quickly run over why we use the terms *sum* and *product* to describe the data types, and show some examples in both Haskell and Java.

Haskell

```

1 data SomeType = A Bool Bool Bool
2               | B Bool
3               | C

```

Here the type `SomeType` has 3 constructors, `A`, `B`, and `C`, where `A` takes 3 parameters, `B` takes one, and `C` zero. The type of `SomeType` could be expressed algebraically as

$$\underbrace{(\text{Bool} \times \text{Bool} \times \text{Bool})}_A + \underbrace{\text{Bool}}_B + \underbrace{1}_C$$

The `Bool` type can take on 2 different values (`False` and `True`), so the constructor can construct $2 * 2 * 2 = 8$ different values, since there are 8 different combinations you can make from 3 booleans (e.g. `A True False False` is one example). The constructor `B` can produce 2 different values, and `C` can only produce one (not zero!). Thus, the total number of values of type `SomeType` is $8 + 2 + 1 = 11$, as the data type is the sum of the three products we've just described.

In short, a sum type denotes “one of” its constituent types (if your function takes `SomeType` as input, it will get either `(A b1 b2 b3)`, `(B b1)` or `C` for some boolean values `b1 . . .`), while a product type denotes “all of” its constituent types (e.g. a value constructed with `A` will have all 3 booleans present.) Another way to express a product type in Haskell is with tuples, e.g.

```

1 type MyTriple = (Bool, Int, Char)

```

Java

In Java, we can model the same kind of data types using classes and inheritance. The instance variables of a class determine a “product” type, e.g.

```

1 class SomeClass {
2     boolean a;
3     boolean b;
4     boolean c;
5 }

```

Similar to the constructor for `A` in the previous section, there are $2 * 2 * 2 = 8$ different values an object of class `SomeClass` can have. Add another boolean, and you get 16 different values. Technically, a variable of type `SomeClass` can take one $8 + 1$ different values, since `null` is also a valid value for all object variables in Java, but sometimes we ignore this fact and tell the users of our functions to kindly not pass in `null` as an argument where we expect an actual object. Now, to get sum types, we might use class hierarchies in Java:

```
1 interface SomeType {}
2
3 class A implements SomeType {
4     boolean a;
5     boolean b;
6     boolean c ;
7 }
8 class B implements SomeType {
9     boolean a;
10 }
11 class C implements SomeType {
12 }
```

Now, an object of type `SomeType` can (ignoring null) has $8 + 2 + 1$ different values, just like in the Haskell example above.

Chapter 2

Anatomy of an Interpreter

An **Interpreter** is a computer program that directly executes instructions written in a programming language. This differs from a **Compilers** which translates a program from one language to another(usually to a lower level one ex. C or ASM).

We usually divide the compiler/interpreter process into two categories, front end, and back end. The front end is the part of the compiler/interpreter that takes the source code and converts it into some intermediate representation. In compilers, this is usually some form of byte code or 3-word code, that can then be translated into the target language. This is not necessary for an interpreter where the intermediary representation is often in the form of an annotated AST. The back end is the part of the compiler/interpreter that takes the intermediary representation and executes it(in the case of interpreters) or translates it back into code(compilers).

2.1 Phases of an interpreter

An interpreter is composed of several phases.

The front end consists of the lexical analyzer, the syntax analyzer, and the semantic analyzer.

The backend is the evaluator.

The **Lexical Analysis** takes the actual characters that the code is made up of and divides it up into its lexical tokens(by the tokenizer) using the concrete syntax of the program¹. Take for instance the following expression:

$$(1 + 2) * 13$$

The lexical analyzer would divide this into the following tokens: ["(", "1", "+", "2", ") ", "*", "13"]

¹Not covered by this course, so you can safely ignore how this works.... *for now!*

This list of tokens is then sent to the **Syntax Analyzer**. The syntax analyzer takes the token list and builds a parse tree out of the tokens and the concrete syntax (fig.2.1).

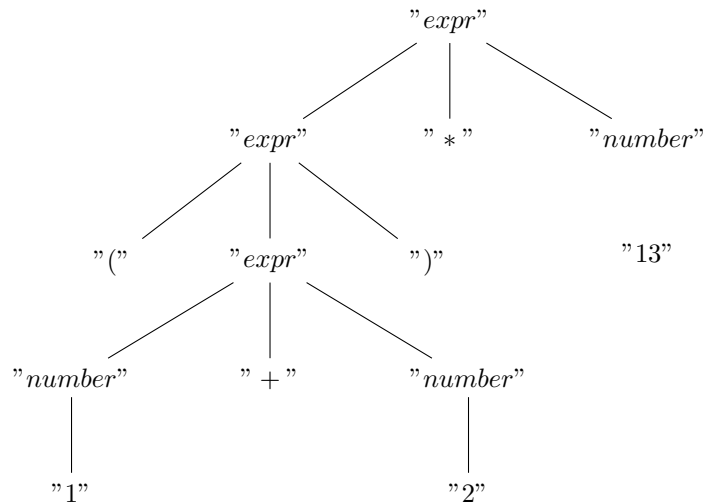


Figure 2.1: Parse tree

The parse tree is then converted into an **Abstract Syntax Tree** by the abstract syntax rules (fig.2.1).




Figure 2.2: AST

The syntax analyzer builds a parse tree out of the tokens and the concrete syntax. This is then converted into an AST by the abstract syntax rules. Which is then handed over to the next part, the **Semantic Analyzer**. The AST is type-checked, checked for well-formedness, names are resolved, and types are inferred. The new AST, now with added information, is then given to the evaluator so that it can be evaluated and produce a result. See Fig. 2.3.




Figure 2.3: Phases of an interpreter

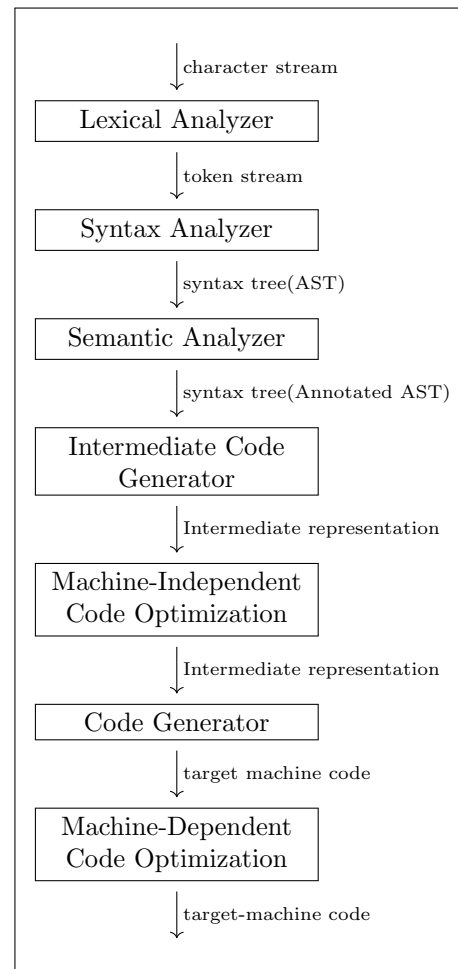


Figure 2.4: Phases of a "normal" compiler

Figure 2.5: Please appreciate the figures above, they were hard to make

2.2 ASTs

An **Abstract Syntax Tree** is the tree representation of the syntactic structure of a program; The abstract syntax describes the structure of the abstract syntax tree - it can be defined using a regular tree grammar, or an algebraic data type or term (in Rascal, ML, Prolog, ...), or an object-oriented inheritance hierarchy of node classes (Java, C++, ...), or as an S-expression (in Lisp languages). The abstract syntax tree can be used as an internal representation in a language processor, but it is not the only possible representation.

An abstract syntax can be generated by a grammar in the following way: For every non-terminal type, there is a corresponding abstract syntax type. Each type has one constructor (or node type) corresponding to each production in the grammar, with one child for every symbol in the production that is not a literal token (e.g., punctuation, keywords, or spaces). If a constructor has only one child, of the same type, it can be removed (e.g., this would be the case for a parenthesis expression). You can do this process entirely based on the information contained in a parse tree.

Translating a parse tree into a corresponding abstract syntax tree is called imploding the parse tree. Given an abstract syntax tree, it is possible to reconstruct a parse tree or program text, given the original grammar - though the resulting program may be slightly different in terms of spaces and punctuation. This is called unparsing or pretty-printing (particularly if the output is nicely formatted). Parsing, imploding, pretty-printing, and then reparsing may not yield the exact same parse tree as the original tree, but it should still implode to the same abstract syntax tree (otherwise there is a bug in your toolchain!).

Various phases in a language processor may change the abstract syntax tree, or use slightly different versions of the abstract syntax (e.g., after type checking, the nodes for variables include the type of the variable) - it is also possible to decorate or annotate the AST as processing proceeds. This adds extra information to the nodes in the AST, without impacting the structure of the abstract syntax. Important abstract syntax design considerations are;

- **Simplicity.** Generally, your compiler tools will do a lot of work on AST, and the fewer different cases you have to worry about, the better. For example, if the processing of overloaded functions and operators is the same (which it is to some degree in C++), you may want to have only one AST node type to cover both. Having a lot of unnecessary nodes in the tree can be annoying as well, and may make processing slower.
- **Good correspondence with the constructs of the language.**
- **Availability of information during processing.** Some information can be computed from the tree (such as type information) and might be encoded directly in the tree (at least at later stages) for easy processing.
- **Being end-user friendly or familiar to most programmers isn't an important consideration** - the abstract syntax may be radically different from the surface concrete syntax if that helps the compiler writer.

2.3 Static Analysis

2.3.1 Type Checking

Programming language typing always falls into one or two categories; static, and dynamic typing. Static typing is when type checking happens at compile-time, while dynamic typing happens during runtime. We will focus on static typing. Statically typed languages typically have these properties.

- Variables and data structures must be declared before use.
- Variables and data structure fields can only hold values of the declared type.
- Operations(i.e functions, procedures, methods) and types must be declared.
- Declaring the exact types of variables and operations isn't always needed. Some languages use type inference (Discussed in 3.2.3).

What is a type checker and how does it work? A type checker is a meta-program that checks that verifies that the type of some construct(lists, expressions, etc.) matches what's expected of it. For example, a type checker will check that the Plus **Expression** takes two Integers.

This lets the type checker discover and report certain errors before the program runs. To do this a type checker needs to know;

- How the language should look.
- The language types.
- Rules for assigning types to the constructs.

Let's do a practical example

Here's the abstract syntax for our Simple Typed Language or STL for short.

```

1 type VarDecl = [(String, ExprType)]
2
3 data Expr =
4     Var String
5     | I Int
6     | B Bool
7     | BinOp Op Expr Expr
8     | UnOp Op Expr
9     | Choice Expr Expr Expr deriving (Show, Eq)
10 data Op = Plus | Mult | Or | And | Not | Eq deriving (Show, Eq)

```

Before we can continue we need to define the types that are allowed. We decide to use two types, Integers, and Booleans.

```

1 data ExprType = Integer | Boolean deriving (Show, Eq)

```

Now to the meat of the exercise, the type checker itself. The usual way to do this in Haskell (and most other languages I've experienced) is to define a series of recursive functions, one for each expression/operand.

```

1 typeCheck :: VarDecl -> Expr -> ExprType

```

I have found it easiest, to begin with, the cases that form the basic "building blocks" of a language, the literals. It is easy to know the type of Int Literals (Integer obviously), and Bool Literals (Boolean).

```

1 typeCheck _ (I _) = Integer
2 typeCheck _ (B _) = Boolean

```

After these "base" cases we add a case for Vars, this is slightly more complicated since the type is dependent on the list of var declarations, so we need to check the VarDecl list.

```

1 typeCheck vars (Var name) =
2     case lookup name vars of
3         Just tp -> tp
4         Nothing -> error $ "No variable could be found named "++name

```

We use a pretty clever Haskell expression called "case". This lets us pattern match the result of the function call in lookup. I strongly recommend that you all learn how to use these since they are extremely useful and I'll be using them liberally during this example. We now move on to the ops.

```

1 typeCheck vars (UnOp Not expr) =
2     case typeCheck vars expr of
3         Boolean-> Boolean
4         _-> error "Argument not boolean"
5
6 typeCheck vars (BinOp Plus left right) =
7     case (typeCheck vars left, typeCheck vars right) of
8         (Integer, Integer)->Integer
9         _->error "One of the args is not an Integer"
10
11 typeCheck vars (BinOp Mult left right) =
12     case (typeCheck vars left, typeCheck vars right) of
13         (Integer, Integer)->Integer
14         _->error "One of the args is not an Integer"
15
16 typeCheck vars (BinOp Or left right) =
17     case (typeCheck vars left, typeCheck vars right) of
18         (Boolean, Boolean)->Boolean
19         _->error "One of the args is not a Boolean"
20
21 typeCheck vars (BinOp And left right) =
22     case (typeCheck vars left, typeCheck vars right) of
23         (Boolean, Boolean)->Boolean
24         _->error "One of the args is not a Boolean"
25
26 typeCheck vars (BinOp Eq left right) = Boolean

```

These all more-or-less follow the same pattern ². They all check that the arguments are of the correct type and return the operand type if so, if not they raise an error. Eq is the odd one out since it returns a boolean no matter what since it checks if two expressions are the same. The last case is the most complex. Choice tests a boolean condition and returns one of the two expressions depending on the value. The problem is that we don't know which branch will return. We have therefore decided both branches need to be of the same type.

```

1 typeCheck vars (Choice test left right) =
2     case typeCheck vars test of
3         Boolean | l == r -> r
4         | otherwise -> error "Args did not match"
5         _-> error "Test condition is not a Boolean"
6
7     where
8         l = typeCheck vars left
9         r = typeCheck vars right

```

Here we begin by checking that the test evals to a boolean type, if so we check that both branches have the same type and return that type. Note that even though an evaluator would only evaluate one of the two branches, we still type-check them both.

²Maya made me use fancier words, apparently "pretty similar" isn't good enough.

2.3.2 Wellformedness

For a program to be **Wellformed** it needs to satisfy all the constraints(think rules) on it. This means that the program follows all the rules for it like;

- The program conforms to the AST.
- It is typed correctly
- All procedure calls/declarations are wellformed.
- and much more.

To check if a program is wellformed we usually implement a so-called constraint checker. These are pretty much just unit tests. The recipe for a constrain checker is as follows;

- **Negative test cases** - Designate one negative test case for each constraint that should be checked. Ideally, each such test case should violate only one constraint.
- **Reporting** - Choose an approach to "reporting". The result of constraint violation may be communicated either through a boolean value, as a list of errors, or by throwing an exception.
- **Modularity** - Implement each constraint in a separate function, thereby allowing modularity and testing.
- **Testing** - The constraint violations must be correctly detected for the negative test cases. The positive test case must pass.

Chapter 3

Store, State, and Storage

3.1 State

3.1.1 Store

When running a program we often want to remember intermediary values or have variables. For us to do this we need some way of storing values. To do this we create a **Store**. A store is very simply an array, usually containing either bytes or ValueTypes. The store is the program's memory (this is true in C). To access a value located at i in our store we simply access the value at array index i . Worth noting that the store is what we call the program heap and by tradition, it grows upwards. This means that new entries are stored at the highest available index in our store. This is because the stack grows upwards and gives us the best possible use of the memory.

3.1.2 Variables and Environment

We now have a place to store stuff, but how do we know where it is in the store? This is where **Environments** comes into play.

An environment is a map between variables and their location in the store.

Elaborate more about variables as name bindings.

3.2 Scoping

Most variables are not visible to the entire program. In previous courses, you have in all likelihood encountered global and local variables when programming. Where global variables are variables that are visible to the entire program, and local variables are variables that are only visible to the method or function that they were declared in.

The term for where a variable is visible is called the **Scope** of a variable. Variable scoping is useful because it lets us keep variables in different parts of our programs separate. If you were to write a calculator it would be somewhat awkward if you could only use x and y once.

Scope can also apply to more than just variables, and all declarations usually have scope. Most languages generally use one of three classes of scoping, and these can have wildly different effects on a program's semantics:

- **Runtime Scoping** - In Runtime Scoping the variable that is in scope is determined by the execution of the program and is the variable that was last seen. Runtime scoping makes no difference between declaration and usage. Variables are declared, initialized, and updated as the interpreter proceeds with the code. Depending on the code's branching structure, a variable may or may not be declared. This means that a variable's scope is often the entire program.
- **Static Scoping** - In Static Scoping the variable that is in scope is determined by the structure of the program. This means that the variable that is in scope is the variable that was declared in the closest enclosing scope. A variable's scope is often just the block where it was declared. Static scoping is easy to reason about and is the most common form of scoping used today.
- **Dynamic Scoping** - In dynamic scoping the scope of a variable is determined by the usage context, instead of its declaration. Dynamic scoping is often hard to reason about since a variable has to be reasoned about in every usage context. Dynamic scoping is therefore not used very often.

To better illustrate the differences let's take a look at an example!

```
1  // A program to demonstrate static scoping.
2
3  int x = 10;
4
5  // Called by g()
6  int f()
7  {
8      return x;
9  }
10
11 // g() has its own variable
12 // named as x and calls f()
13 int g()
14 {
15     int x = 20;
16     return f();
17 }
18
19 int main()
20 {
21     print(g());
22     return 0;
23 }
```

Figure 3.1: Scoping Example

Runtime Scoping

If we execute the program with runtime-scoping the program will return 20. This is because the variable `x` is first declared in line 3, but is then overridden in `g()` on line 5. When we call `g()` on line 9, the variable `x` is set to 20, and this is then returned by `f()`.

Static Scoping

If we execute the program with static scoping the program will return 10. This is because the variable `x` is first declared on line 3, before being shadowed in `g()`, but because `x` is only 20 within `g()`, the `x` in `f()` is still the same as it was on line 3.

Dynamic Scoping

The result of executing the program with dynamic scoping is 20 because the variable `x` is set to 20 in `g()` and `g()` calls `f()`.

3.3 Data Structures in Memory

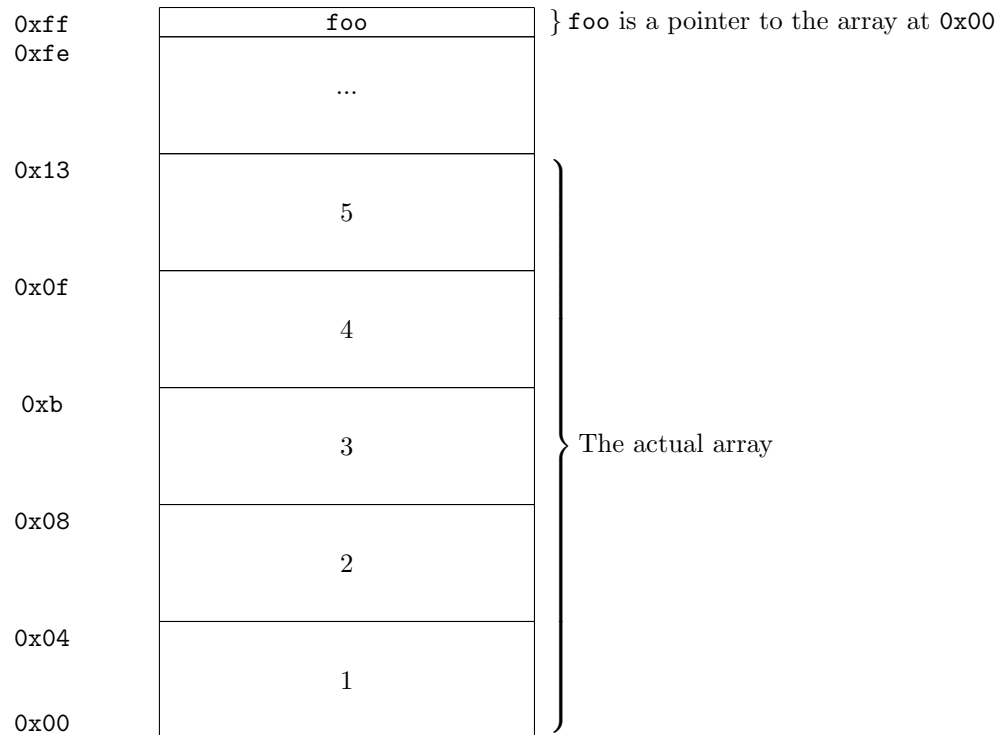
3.3.1 Arrays

An array is a type of data structure that holds a fixed number of elements of the same type. Arrays are usually stored in memory as a contiguous block of memory. Arrays are usually indexed by integers.

```
1  int foo[5] = {1, 2, 3, 4, 5};
```

Figure 3.2: Array of Ints with 5 elements.

In the code above we have an array consisting of 5 integers, we'll assume that addresses only use 1 byte of memory, and integers take up 4 bytes. Figure 3.3 shows how the array would be stored in memory.

Figure 3.3: `foo` in memory

The array is stored as a contiguous block of memory, and the variable `foo` is a pointer on the stack that points to the start of the array. To access the elements in the array we can use the pointer `foo` and add an offset to it. The address for the n 'th element in an array is given by the following formula, where

$$\text{foo}[n] = \&\text{foo} + n * \text{sizeof}(\text{int})$$

- `&foo` is the pointer to the array
- `n` is the index of the element we want to access
- `sizeof(int)` is the size of an integer in bytes

3.3.2 Matrixes and Multidimensional Arrays

Things get a little bit more complicated when dealing with multidimensional arrays, and the exact nature of how multidimensional arrays are stored in memory depends on the language. There are two main ways of storing arrays of arrays. First is to that each entry in the outer array is a pointer to the nested array. This has a few problems, the first is that the cpu needs to do an extra memory lookup to get the address of the nested array. The second is that it becomes impossible for the compiler to make memory layout optimizations that would otherwise be possible.

The second way is to store the entire array as a contiguous block of memory. This has the advantage that the cpu can do a single memory lookup to get the address of the nested array,

and the compiler can make optimizations. The disadvantage is that the memory layout is not as intuitive(See Figure 3.4).

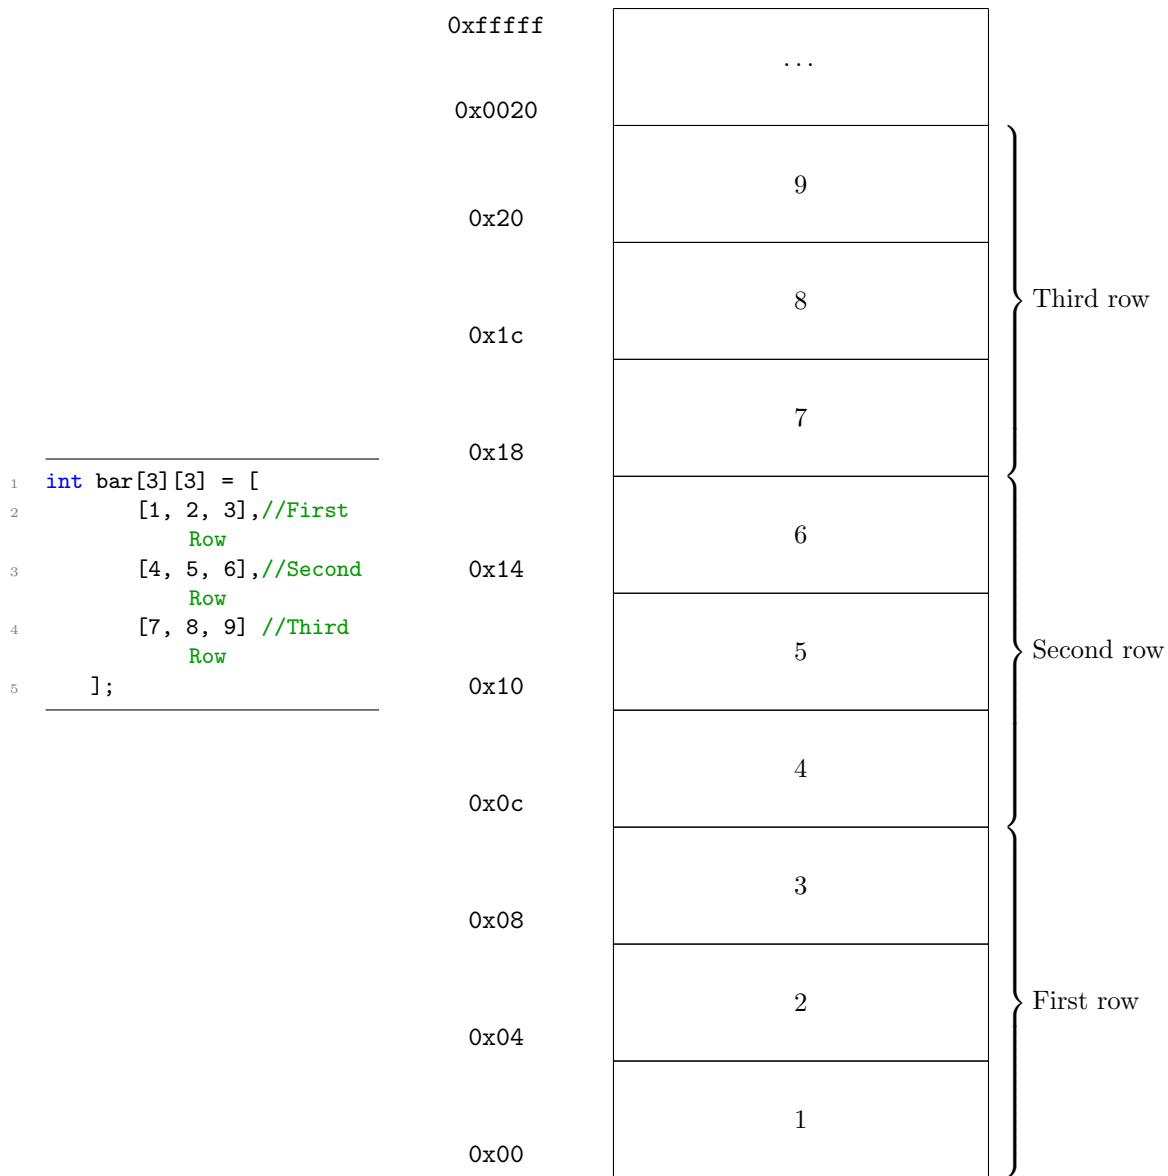


Figure 3.4: `bar` in memory with contiguous layout

3.3.3 Records, Structs

Records are a data structure that can hold multiple values of different types. Records are also known as structs, tuples, or classes depending on the language.

```

1 struct foo {
2     int x; //Offset 0
3     bool y; //Offset 4
4     double z; //Offset 5
5 };

```

$$\begin{aligned}
 |\text{foo}| &= |x| + |y| + |z| \\
 &= 4 + 1 + 8 = 13
 \end{aligned}$$




Figure 3.5: foo in memory

Arrays of Records

Storing records in memory is very straight forwards, as with arrays, there are two main ways of storing arrays of records. The first is the one used by Java, where each element of the array is really a pointer to a record thats stored somewhere else on the heap.

The other way is to store them as you would any other datatype, as a contiguous block of memory, with one record stored after the other.

To access a field in a record we can use the following formula, where

$$\text{arr}[n].x = \&\text{arr} + n * |\text{foo}| + |\text{offset}(x)|$$

where

- `arr` is the Array
- `n` is the index of the element we want to access
- `foo` is the struct or record
- `x` is the field we want to access

Chapter 4

Procedures

A procedure is a programming construct that abstracts away the implementation of an algorithm. Instead of writing 20 lines of code every time you want to run an algorithm we simply call a procedure with those 20 lines of code inside it. This increases the readability of our code and makes it easier to write it.

4.1 Procedure Declarations

In this course, we have a simplified procedure declaration inspired by PASCAL. Our procedure declarations have a list of parameters(variables that pass in/out of the procedure), another list of local variables(variables that are used within), and a single statement(the algorithm itself). If we were to translate this abstract syntax into something more concrete then a typical procedure might look something like this.

```
1 procedure p (upd a: integer , out b: integer ; obs c : boolean ) ;
2 var x: integer ;
3   y: boolean ;
4   z : boolean ;
5 begin
6   if c then begin b := a ; x := b ; a := x end;
7   y := not c ;
8   c := y or c ;
9 end;
```

The abstract syntax would be something like this in Haskell;

```
1 -- | Procedure declaration
2 data ProcedureDeclaration = Proc
3     String -- Name of the procedure
4     [Parameter] -- Parameter list
5     [VarDecl] -- Local variables
6     Stmt -- Statement part
7     deriving (Show, Eq, Read)
8
9 -- | Procedure parameters: mode and variable declaration
10 type Parameter = (Mode, VarDecl)
11
```

```

12 -- | Parameter modes: observe, update, output
13 data Mode = Obs | Upd | Out
14     deriving (Show, Eq, Read)
15
16 -- | Variable declaration: variable name and its typ
17 type VarDecl = (Var, Type)

```

4.1.1 Parameters & Local Variables

A **Local Variable** is a variable that only exists in a specific part of the program. We're going to talk about them in the context of procedures. These local variables are declared or defined and only used within that procedure. Since the procedure exists in a vacuum, the local variable is allowed to have the same name as variables in other parts of the program. Since the variable is stored at a different store location any change to the local variable wouldn't change the variable outside.

Parameter are the variables the procedure uses to communicate to the outside world. They are variables, but with one key difference; They also specify *how* it communicates with the outside world, the parameters specify if the variable is an output variable(write), observed only(read), or updatable(is this a word?)(read and write).

Another way of looking at parameters is that parameters are local variables that are initialized with the passed arguments at invocation time.

4.1.2 Performing a Procedure

When performing a function we can usually assume that the parameters have already been added to the environment and initialized and that except for those parameters we have a clean environment to work with. The performing of a procedure consists of the following steps;

1. The first is that we need to somehow remember the current environment or delete all the local variables when we're done. The best way is to get the current **Stackframe**.
2. We then add all the local variables.
3. We then execute the procedure statements.
4. We then reset the environment back to how it looked before using the saved stack frame from (1). This ensures that all local variables that shouldn't exist outside the procedure are removed.

4.2 Executing a Procedure Call

We now know how to actually perform a procedure, but for us to even be able to perform a procedure in the first place we need to prepare the environment and program for it. There are two main ways of making the program ready to perform a procedure. These are dependent on what type of **Argument** passing we are doing.

4.3 Arg Passing

Argument passing falls into one of two camps.

4.3.1 Copy Semantics

The first is **Copy Semantics**. In copy semantics, all the arguments are passed to the parameters by copying the value into them. With copy semantics, the procedure arguments are first evaluated and then bound to the parameters when they are added. The procedure is then performed, and the results are obtained. Those are then copied back to all the arguments that are upd or out. To further explain let's look at a trivial example to understand what happens in-store. Say we have a procedure that simply adds 5 to any number.

```

1 procedure add5(obs x: integer, out res: integer)
2 begin
3     res = x + 5;
4 end;
```

Now let's look at the main procedure

```

1 procedure main()
2     var a: integer;
3     var b: integer;
4 begin
5     a = 42;
6     add5(a, b);
7 end;
```

Figure 4.1 shows us what the state looks like before we start executing the procedure call. The


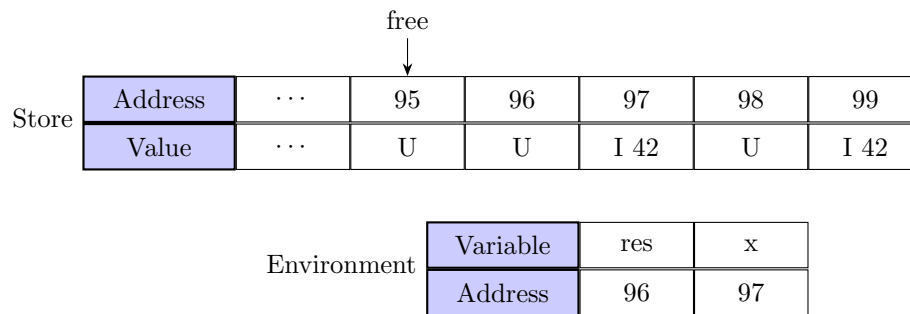
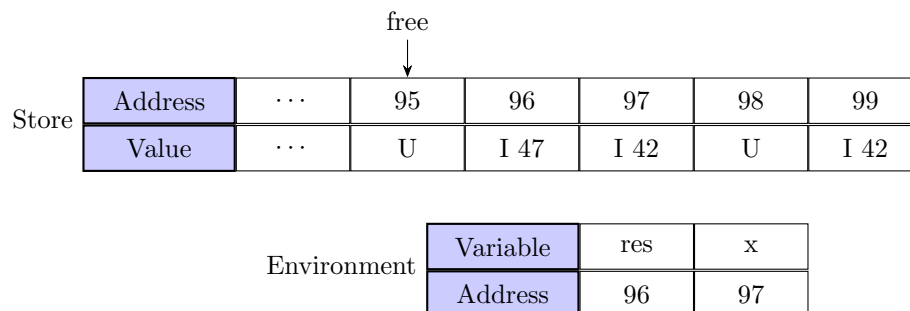


Figure 4.1: State before the procedure call

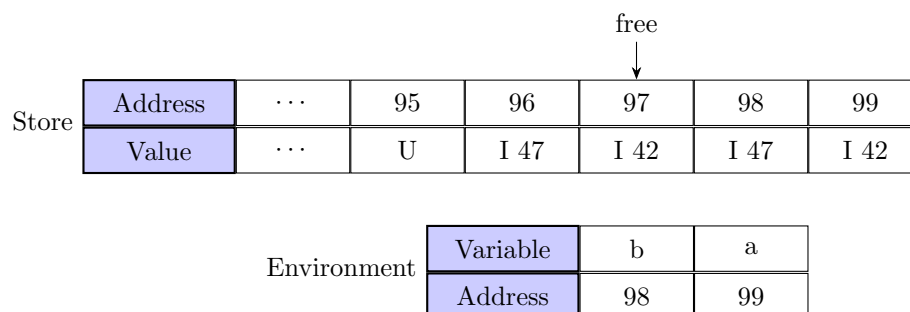
first thing we do when entering the procedure call is to save the stackframe so we can restore the environment to its proper place once we're done. We then clear the environment, ensuring that the only variable declared by the procedure is in scope. We then add all the parameters (in our case we add x and res) and copy the values of a and b into x and res respectively ¹. Figure 4.2 shows what the state looks like at this point.

We have now prepared the procedure so that it can be performed. Figure 4.3 shows how the state looks after its been performed.

¹since res is out it isn't initialized

Figure 4.2: State when entering `add5` (copy semantics).Figure 4.3: State after `res = x + 5` in `add5` (copy semantics).

Now comes the fun part, *cleaning*! Just performing the function isn't enough, if we were to just drop it here we would have an environment that is drastically different than then when we started. We could just reset the stackframe to what it was when we started, but then `b` wouldn't get its new value. What we need to do is copy the values of the out/upd params back to their argument variables. Then we can reset the stackframe making our state look like fig. 4.1. As you can see the `free` pointer now points to where `x` used to point and `b` now has the same value as `a`. That means that the next time we add something to the environment it will overwrite those values in the store.

Figure 4.4: State after `add5(a, b)` in `main` (copy semantics).

4.3.2 Reference Semantics

Now, what if `add5` could instead write *directly* to `b` in `main` instead of having to allocate its own copy which is later copied? This is the main advantage of **Reference Semantics**. Lets reuse the example from before. As you can see the from fig. 4.5 the state before we enter the procedure call is the same as during copy semantics.




Figure 4.5: State before the procedure call

It is now we encounter our first change. Like before we get the stackframe before clearing the environment and adding the parameters, but here comes the change. instead of allocating space and copying the value of the args to the upd/out params we instead set their address to the address of the corresponding argument. We now get a state that looks like fig. 4.6. Since `x` is an obs parameter it is still allocated and copied to as normal.

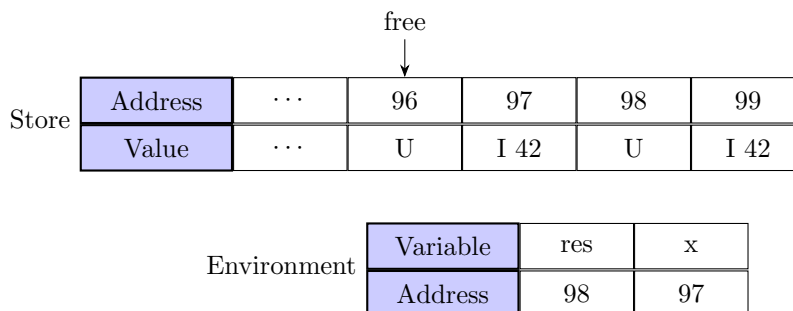



Figure 4.6: State when entering `add5` (reference out).

Since `res` now points to the same place as `b`, any changes to `ref` will also be reflected by `b`. This makes cleanup much easier! infact all we have to do after performing the procedure is to reset the enviroment back to where it was and call it a day makng the state look like fig. 4.7.

Figure 4.7: State after $\text{res} = x+5$ in `add5` (reference out).

Chapter 5

Signatures

So far we've defined all our operations as part of our AST. By introducing signatures into our compilers we can abstract away the implementation of our operations from the AST. This allows us to change the implementation of our operations without changing the AST. Before we can do this we need to talk a bit about the theory behind signatures and algebras

5.1 A little bit of Theory

5.1.1 Signatures

Signatures are a way to define a set of operations and their types.

Formaly a signature¹ is defined as $I = \langle S, F \rangle$ where

- S is a set of sorts(aka typenames), and
- F is a set of function declarations $f : s_1, \dots, s_n \rightarrow s$, for $s_1, \dots, s_n, s \in S$

Signatures alone do not define any semantics, they are just a way to define which operations and types exist, but not what they do. To define the semantics we need to define what we call an algebra.

Here's an example of a signature.

$$Nat = \langle \{N\}, \\ \{zero : N, \\ succ : N \rightarrow N\} \rangle$$

5.1.2 Algebras

An **Algebra** is a way to define the semantics of a signature. An algebra assigns every sort to a domain and every function to a function on the domains. An algebra A for a signature $I = \langle S, F \rangle$ defines

- a set $\llbracket s \rrbracket_A$ for every sort $s \in S$, and
- a total function $\llbracket f \rrbracket_A : \llbracket s_1 \rrbracket_A \times \dots \times \llbracket s_k \rrbracket_A \rightarrow \llbracket s \rrbracket_A$ for every $(f : s_1, \dots, s_k \rightarrow s) \in F$

¹also called an *Interface*

It's possible to have multiple algebras for a signature. Here are a couple of algebras for the *Nat* signature.

$$\begin{aligned}\llbracket N \rrbracket_{A_1} &= \mathcal{N} \\ \llbracket zero \rrbracket_{A_1} &= 0 \\ \llbracket succ \rrbracket_{A_1} &= \lambda n \mapsto n + 1\end{aligned}$$

$$\begin{aligned}\llbracket N \rrbracket_{A_2} &= \{1, 2, 7\} \\ \llbracket zero \rrbracket_{A_2} &= 2 \\ \llbracket succ \rrbracket_{A_2} &= \lambda n \mapsto 7\end{aligned}$$

5.2 Implementing Algebraic Specifications

Integrating signatures and algebras into our interpreter means we can use Algebraic specification theory to formalize and reason about it. It also lets us change the implementation of our operations without changing the AST and is one step towards user-defined types, and generic programming.

Lets take a look at how we can rewrite an interpreter to make use of signatures.

```

1  type Env = [(String, Int)]
2
3  data Expr
4    = Lit Int
5    | Add Expr Expr
6    | Sub Expr Expr
7    | Mult Expr Expr
8    | Div Expr Expr
9    | LEQ Expr Expr
10   | Var String
11   deriving (Show, Eq)
12
13 eval :: Env -> Expr -> Int
14 eval _ (Lit n) = n
15 eval env (Add e1 e2) = eval env e1 + eval env e2
16 eval env (Sub e1 e2) = eval env e1 - eval env e2
17 eval env (Mult e1 e2) = eval env e1 * eval env e2
18 eval env (Div e1 e2) = if eval env e2 /= 0 then div (eval env e1) (eval env e2) else
19   error $ "Division by zero: " ++ show e1 ++ " / " ++ show e2
20 eval env (LEQ e1 e2) = if eval env e1 <= eval env e2 then 1 else 0
21 eval env (Var string) = case lookup string env of
22   Just n -> n
23   Nothing -> error $ "Variable " ++ string ++ " not found in
24     environment"

```

Figure 5.1: Interpreter for a simple language

We do this by creating a new AST without any operations apart from Literals, Variables, and Function calls. And instead of returning a `Int` we make it so that the AST can work for any `valuedomain`(Figure 5.2).

```

1 type Env valueDomain = [(String, valueDomain)]
2 data Expr valueDomain
3     = Lit valueDomain
4     | Var String
5     | FunCall String [Expr valueDomain]

```

We also change the evaluator to reflect this change. One major change(apart from the lack of operations) is that we now need to pass the algebra(`funmod`) to the evaluator.

```

1 eval :: (String -> [valueDomain] -> valueDomain) -> Env valueDomain ->Expr valueDomain
   -> valueDomain
2 eval fmod env (Lit n) = n
3 eval fmod env (Var string) = case lookup string env of
4     Just n -> n
5     Nothing -> error $ "Variable " ++ string ++ " not found in
   environment"
6 eval fmod env (FunCall f args) = fmod f (map (eval fmod env) args)

```

Not that we've adapted the AST for signatures it's time to implement a signature for the operations we removed.

```

1 intrinsics :: Signature
2 intrinsics = (["Int", "Bool"],[
3     ("add", ["Int", "Int"], "Int"),
4     ("sub", ["Int", "Int"], "Int"),
5     ("mult", ["Int", "Int"], "Int"),
6     ("div", ["Int", "Int"], "Int"),
7     ("leq", ["Int", "Int"], "Bool")
8 ])

```

You may have noticed one major change from the old AST. The old AST only worked on `Ints`, but we've added a second sort `Bool` to the signature. Since Haskell only lets us use one type for the `valuedomain` we need to create a new type that can hold both `Ints` and `Bools`.

```

1 data VD = Bool Bool | Int Int

```

We then define our algebra for the functions in the signature.

```

1 intrinsicSemantics :: String -> [VD] -> VD
2 intrinsicSemantics "add" [Int a, Int b] = Int (a + b)
3 intrinsicSemantics "sub" [Int a, Int b] = Int (a - b)
4 intrinsicSemantics "mult" [Int a, Int b] = Int (a * b)
5 intrinsicSemantics "div" [Int a, Int 0] = error $ "Division by zero: " ++ show a ++ " /
  0"
6 intrinsicSemantics "div" [Int a, Int b] = Int (div a b)
7 intrinsicSemantics "leq" [Int a, Int b] = Bool (a <= b)

```

Now, there are some downsides to using signatures. For one we no longer get to use Haskell's built-in type system to check that we're using the right types. That means that we need to implement and adapt our type checker to work with signatures.

5.2.1 ADTs

We've implemented Signatures in our interpreter, but we haven't given the user a way to define their own yet. We can do this by introducing **Abstract Data Type (ADT)** in our language. Abstract Data Types are a way to define a type and its operations without exposing the implementation.

The best example of ADTs are probably Interfaces in Java.

Methods in Java Interfaces don't have any implementation², they just define the signature of the method. This means that we can define a method in an interface and then implement it in multiple ways. This is a very powerful tool for generic programming.

```

1 package ADT;
2
3 public interface IStack<T> {
4
5     public void push(T item);
6
7     public T pop();
8
9     public T peek();
10
11 }

```

Figure 5.2: ADT for a Stack

²We'll conveniently ignore the existence of the `default` keyword.

Figure 5.2.1 shows one implementation of `IStack`.

```

1  class Stack<T> implements IStack<T> {
2
3      private List<T> stack;
4      public Stack() {
5          stack = new ArrayList<T>();
6      }
7      public void push(T item) {
8          stack.add(item);
9      }
10
11     public T pop() {
12         return stack.remove(stack.size()-1);
13     }
14
15     public T peek() {
16         return stack.get(stack.size()-1);
17     }
18
19 }

```

Figure 5.3: Stack

Since ADTs don't specify an implementation it is perfectly valid for me to implement whatever I want as long as it has the same signature. For example, I could implement `IStack` as a linked list instead of an array. Or we could go a step further and observe that the only difference between stacks and queues is the behavior of `pop`.

So Figure 5.2.1 is a totally valid implementation of the `Stack` ADT.

```

1  public class Queue<T> implements IStack<T>{
2      private List<T> queue;
3      public Queue() {
4          queue = new ArrayList<T>();
5      }
6
7      public void push(T item) {
8          queue.add(item);
9      }
10
11     public T pop() {
12         return queue.remove(0);
13     }
14
15     public T peek() {
16         return queue.get(0);
17     }
18 }

```

Figure 5.4: Queue

5.3 Generic Programming

5.3.1 Concepts

Appendix A

Assertions

Assertions are a statement that checks that some predicate holds.

Assertions ensure that a program is always in a valid state. If for any reason an assertion fails a program crash is the result. This is usually the preferred option since an invalid state means that the program behaviour is undefined, and can result in security or safety issues.

A.1 Types of Assertions

Pre-Conditions

Pre-conditions are assertions that must hold before a function or procedure is called. They are used to ensure that the function is called with valid arguments and that the function can be executed correctly.

Post-Conditions

Post-conditions are assertions that must hold after a function or procedure is called. They are used to ensure that the function has been executed correctly and that the result is valid.

Invariants

Invariants are assertions that must hold at all times. They are used to ensure that the program is always in a valid state.

For example, we can use assertions to make sure that an implementation of the natural numbers is valid. We could do this by always checking that any argument, value, or result is always greater than or equal to zero.

Appendix B

Language Standards

B.1 Software Engineering Implications of Languages

The choice of language can have a massive impact on software development. The move to higher abstraction languages has caused a corresponding increase in efficiency since the developer can focus all their efforts on what the program should accomplish instead of having to work on the details. **Software Language Engineering** has many applications within software engineering like;

- Design, implementation, and usage of DSLs that are tailor-made for a specific problem or technical domains, like, UI, web services, configurations, testing, data exchange, interoperability, deployment, and distribution.
- Software reverse engineering and re-engineering in many forms, for example, analysis of projects regarding their dependence on open source software, integration of systems, and migration of systems constrained by legislation or technology.
- Data extraction in the context of data mining, information retrieval, machine learning, big data analytics, social science, digital forensics, and AI, with diverse input, and artifacts to be parsed in interchange formats to conform to.

Software languages can also impact the security of any product. If a language has a fundamental flaw or creates unknown pitfalls that might be hard to spot due to the inherent design of the language could then be exploited by malicious actors.

B.2 Reading specifications

B.2.1 Backus-Naur Form

Backus-Naur Form is a format for describing context-free grammar. We use it to describe the syntax of languages. Although it is probably not part of the curriculum it's still important to know since most language specification describes their languages using some variation of BNF most common of which is the Extended Backus-Naur form(EBNF). An EBNF consists of terminal symbols and non-terminal production rules which are the restrictions governing how terminal symbols can be combined into a legal sequence. Examples of terminal symbols include alphanumeric characters, punctuation marks, and whitespace characters. These constructions

often end up looking like Haskell data structures, this should hopefully help you understand them.

```
1 <symbol> ::= __expression__
```

Here is the complete PASCAL-like language that only allows assignments. in its EBNF form.

```
1 (* a simple program syntax in EBNF - Wikipedia *)
2 program = 'PROGRAM', white space, identifier, white space,
3         'BEGIN', white space,
4         { assignment, ";", white space },
5         'END.' ;
6 identifier = alphabetic character, { alphabetic character | digit } ;
7 number = [ "-" ], digit, { digit } ;
8 string = '"', { all characters - '"' }, '"';
9 assignment = identifier , ":", ( number | identifier | string ) ;
10 alphabetic character = "A" | "B" | "C" | "D" | "E" | "F" | "G"
11                      | "H" | "I" | "J" | "K" | "L" | "M" | "N"
12                      | "O" | "P" | "Q" | "R" | "S" | "T" | "U"
13                      | "V" | "W" | "X" | "Y" | "Z" ;
14 digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9" ;
15 white space = ? white space characters ? ;
16 all characters = ? all visible characters ? ;
```

EBNF also includes regex like syntax for expressing repetition(*,+), optionality(?,[]), and alternatives(—). If the above looks confusing or unintuitive I wouldn't worry. The syntax of EBNF is fairly intuitive and most standards have fairly little of it at once.

Glossary

Abstract Data Type(ADT) An ADT is a data type that is defined by its behavior from the point of view of a user, of the data, specifically in terms of possible values, possible operations on data of this type, and the behavior of these operations.. 33

Abstract Syntax Tree An Abstract Syntax Tree(AST) is a tree representation of the syntactic structure of our program. Can be represented by using trees or terms, and described by an algebraic data type or regular tree grammar.. 11, 12

Algebra An algebra defines the semantics of a signature.. 30

Argument An argument is a value provided to the procedure when it is run. When the procedure is run the parameters of the procedure are initialized with its corresponding argument.. 25

Backus-Naur Form Formal notation for describing grammars. Used to describe the syntax of a language.. 37

Basic Imperative Programing Language An extension of BTL, but with statements and control flow.. 6, 7

Basic Typed Language Simple language with a few types and expressions on those types. 6, 7

Compiler A compiler is a program that translates computer code(*the source language*) into another language(*the target language*) Compilers usually convert some high-level language to some lower-level language.. 10

Copy Semantics Type of argument passing where the parameters are initialized with the value of the arguments. 26

Domain Specific Languages A language(i.e. not just a library) with abstractions targeted at a specific problem domain.

- *External DSL* - A DSL is defined as a separate programming language.
- *Internal/Embedded DSL* - A DSL is defined as a language-like interface to a library.

. 6

Enviroment Map describing where things are located in the **Store**. Kinda like a phonebook. 17

Expression An expression is a syntactic construct that can be evaluated in order to obtain its value. The resulting value is usually one of the program's types.. 13

- General Purpose Languages** A language suited for a wide variety of problems and situations but lacks specialized features to deal with specific programs like a DSL.. 7
- Interpreter** An interpreter is a program that directly executes instructions written in a programming language.. 10
- Lexical Analysis** Lexical analysis is the process of converting a sequence of characters into a sequence of tokens.. 10
- Local Variable** A local variable is a variable that only exists within a limited scope.. 25
- Meta-Language** The language a compiler is written in. 7
- Meta-Programming** Metaprogramming is a programming technique in which computer programs treat other programs as their data. It means that a program can be designed to read, generate, analyze, or transform other programs, and even modify itself while running.. 7
- Object Language** The language that gets compiled. 7
- Parameter** A parameter is a local variable that is initialized with the arguments. Also often contains how those args are to be treated.. 25
- Procedural Imperative Programming Language** An extension of BIPL, but with procedures and functions.. 6
- Reference Semantics** Type of argument passing where the parameters point to the address as the argument.. 28
- Scope** A collection of identifier bindings . i.e. what's captured by the environment at some point in the code.. 17
- Semantic Analyzer** Takes an AST and checks if it follows the rules of the language. Outputs an annotated AST. 11
- Semantics** The semantics of a program concerns the meaning of the program. i.e. what it does. It can take many forms, sometimes we're only interested in the result of the program. Other cases concern the steps taken by the program to reach the said output.. 7
- Signature** A signature defines a set of operations and their types, but not the semantics of the operations.. 30
- Software Language Engineering** Software Language Engineering is the scientific field that researches language development, and maintenance of formal descriptions, and tooling of software languages. 37
- Stackframe** The stackframe is a snapshot of how the program environment looks at a certain point in time. The stackframe saves things that could be changed by running the procedure and lets us restore the program to the previous state without all the changes made by the procedure.. 25
- Store** Program memory, byte/value array, grows upwards.. 17, 39
- Sum of Products** See chapter 1.3. 8
- Syntax Analyzer** Takes a stream of tokens and checks if it follows the rules of the grammar. Outputs a parse tree that is then assembled into an AST. 11

Syntax Syntax refers to the rules that define the structure of a language. Syntax in computer programming means the rules that control the structure of the symbols, punctuation, and words of a programming language.. 7

Wellformed Wellformedness is when a program is following all the rules. smileemoji. 16

Bibliography

- [1] Ralf Lämmel. *Software Language Engineering*. Springer International Publishing, 2018.
- [2] Anya Bagge w/Ralf Lämmel Vadim Zaytsev. Inf225 notes. Lecture Notes for INF225, University of Bergen, Norway, 2016.